

Overview of the course

1. Introduction
2. Apache Spark's basic concepts
3. **Spark SQL**
4. Mllib

Data Frame

- A Data Frame is a distributed collection of data, organised as columns with an associated name.
 - The concept is similar to SQL tables
- The entry point to Spark SQL is the SQLContext class:

Python:

```
from pyspark.sql import SQLContext  
sql = SQLContext(sc)
```

Creating a Data Frame

- Data Frames can be loaded from different sources: JSON, JDBC/ODBC and Apache Hive.

- Check official documentation

- In addition, Spark allows creating data frames from RDDs

```
In [13]: rdd = sc.parallelize([("Maria",35),  
("Jose",42), ("Antonia", 25)])
```

```
In [14]: people = sql.createDataFrame(rdd)
```

```
In [15]: people.show()
```

```
+-----+----+  
|      _1| _2|  
+-----+----+  
|  Maria| 35|  
|   Jose| 42|  
|Antonia| 25|  
+-----+----+
```

Providing extra information to Data Frames

```
from pyspark.sql import Row  
from pyspark.sql.types import *  
rdd = sc.parallelize([("Maria",35), ("Jose",42),  
                      ("Antonia", 25)])
```

```
schema = StructType([  
    StructField("name", StringType()),  
    StructField("age", IntegerType())])
```

```
sql = SQLContext(sc)  
people = sql.createDataFrame(rdd,schema)  
people.show()
```

```
+-----+----+  
|   name|age|  
+-----+----+  
|  Maria| 35|  
|   Jose| 42|  
|Antonia| 25|  
+-----+----+
```

--- Common operations and transformations for DF's

- **agg(*expressions)** returns a new DF with a single row, containing the results of the expressions passed by parameter

```
people.agg(avg(people.col("age")), max(people.col("age")), min(people.col("age"))).show()
```

avg(age)	max(age)	min(age)
34.0	42	25

- **corr(col1, col2), cov(col1, col2)** returns the correlation/covariance between two columns
- **drop(col)** returns a new DF with the specified column dropped from the original
- **withColumn(name, expr)** returns a new column, given a name and the expression that provides its value:

```
people.withColumn("female", people.col("name").endsWith("a")).show()
```

name	age	female
Maria	35	true
Jose	42	false
Antonia	25	true

- Other methods common to RDDs: count, distinct, filter, ...

SQL-like operations on DFs

- **select(*cols)** returns a new DF with only the specified columns:

```
people.select("name").show()
```

name
Maria
Jose
Antonia

- **filter(condition), where(condition)** filters the rows given a condition

```
people.where(people.col("age") < 30).show()
```

name	age
Antonia	25

- **groupBy(*columns)** similar to SQL GROUP BY, providing aggregation functions:

```
people.groupBy(people.col("age") % 10).avg().show()
```

(age % 10)	avg(age)
2	42.0
5	30.0

SQL text queries

```
from pyspark.sql.types import *
sql = SQLContext(sc)
schema = StructType([
    StructField("name", StringType()),
    StructField("age", IntegerType()),
    StructField("passport", StringType())])

students = sql.createDataFrame(
    sc.parallelize([
        ("Jaime", 32, "12345-f"),
        ("Maria", 19, "22222-g"),
        ("Alex", 23, "65432-z")]), schema)

students.registerTempTable("students")
```

SQL text queries (II)

```
schema2 = StructType([
    StructField("passport", StringType()),
    StructField("year", IntegerType())])

enrollments = sql.createDataFrame(
    sc.parallelize([
        ("12345-f", 1990),
        ("22222-g", 2014),
        ("65432-z", 2009)]), schema2)

enrollments.registerTempTable("enrollments")
```

SQL text queries (III)

```
sql.sql("""
    SELECT students.name, enrollments.year AS
enrollment_year
    FROM students, enrollments
    WHERE students.passport = enrollments.passport
    ORDER BY students.name ASC
    """).show()
```

name	enrollment_year
Alex	2009
Jaime	1990
Maria	2014

Hands-on: Google Cluster data analyser

<https://github.com/mariomac/patc-spark/tree/master/exercises/2-sql>